



Hardened/Overview of POSIX capabilities

From Gentoo Wiki

< [Hardened \(/wiki/Hardened\)](#)

[Jump to:navigation](#) [Jump to:search](#)

POSIX capabilities are a partitioning of the all powerful root privilege into a set of distinct privileges.

Contents

- 1 CAP_CHOWN
- 2 CAP_DAC_OVERRIDE
- 3 CAP_DAC_READ_SEARCH
- 4 CAP_FOWNER
- 5 CAP_FSETID
- 6 CAP_FS_MASK
- 7 CAP_KILL
- 8 CAP_SETGID
- 9 CAP_SETUID
- 10 CAP_SETPCAP
- 11 CAP_LINUX_IMMUTABLE
- 12 CAP_NET_BIND_SERVICE
- 13 CAP_NET_BROADCAST
- 14 CAP_NET_ADMIN
- 15 CAP_NET_RAW
- 16 CAP_IPC_LOCK
- 17 CAP_IPC_OWNER
- 18 CAP_SYS_MODULE
- 19 CAP_SYS_RAWIO
- 20 CAP_SYS_CHROOT
- 21 CAP_SYS_PTRACE
- 22 CAP_SYS_PACCT
- 23 CAP_SYS_ADMIN
- 24 CAP_SYS_BOOT
- 25 CAP_SYS_NICE
- 26 CAP_SYS_RESOURCE
- 27 CAP_SYS_TIME
- 28 CAP_SYS_TTY_CONFIG
- 29 CAP_MKNOD
- 30 CAP_LEASE

CAP_CHOWN

CODE CAP_CHOWN

CAP_CHOWN

In a system with the `[_POSIX_CHOWN_RESTRICTED]` option defined, this overrides the restriction of changing file ownership and group ownership.

CAP_DAC_OVERRIDE

CODE CAP_DAC_OVERRIDE

CAP_DAC_OVERRIDE

Override all DAC access, including ACL execute access if `[_POSIX_ACL]` is defined. Excluding DAC access covered by `CAP_LINUX_IMMUTABLE`.

CAP_DAC_READ_SEARCH

CODE CAP_DAC_READ_SEARCH

CAP_DAC_READ_SEARCH

Overrides all DAC restrictions, regarding read and search on files and directories, including ACL restrictions, if `[_POSIX_ACL]` is defined. Excluding DAC access covered by `CAP_LINUX_IMMUTABLE`.

CAP_FOWNER

CODE CAP_FOWNER

CAP_FOWNER

Overrides all restrictions about allowed operations on files, where file owner ID must be equal to the user ID, except where `CAP_FSETID` is applicable. It doesn't override MAC and DAC restrictions.

CAP_FSETID

CODE CAP_FSETID

CAP_FSETID

Overrides the following restrictions, that the effective user ID shall match the file owner ID, when setting the `S_ISUID` and `S_ISGID` bits on that file; that the effective group ID (or one of the supplementary group IDs) shall match the file owner ID when setting the `S_ISGID` bit on that file; that the `S_ISUID` and `S_ISGID` bits are cleared on successful return from `chown(2)` (not implemented).

CAP_FS_MASK

CODE **CAP_FS_MASK**

```
CAP_FS_MASK
    Used to decide between falling back on the old suser() or fsuser().
```

CAP_KILL

CODE **CAP_KILL**

```
CAP_KILL
    Overrides the restriction, that the real or effective user ID of a process,
    sending a signal, must match the real or effective user ID of the process,
    receiving the signal.
```

CAP_SETGID

CODE **CAP_SETGID**

```
CAP_SETGID
    Allows setgid(2) manipulation;
    Allows setgroups(2);
    Allows forged gids on socket credentials passing.
```

CAP_SETUID

CODE **CAP_SETUID**

```
CAP_SETUID
    Allows set*uid(2) manipulation (including fsuid);
    Allows forged pids on socket credentials passing.
```

CAP_SETPCAP

CODE **CAP_SETPCAP**

```
CAP_SETPCAP
    Transfer any capability in your permitted set to any pid, remove any capabili
ty in your permitted set from any pid.
```

CAP_LINUX_IMMUTABLE

CODE **CAP_LINUX_IMMUTABLE**

```
CAP_LINUX_IMMUTABLE
    Allow modification of S_IMMUTABLE and S_APPEND file attributes.
```

CAP_NET_BIND_SERVICE

CODE **CAP_NET_BIND_SERVICE**

```
CAP_NET_BIND_SERVICE
```

```
    Allows binding to TCP/UDP sockets below 1024;
```

```
    Allows binding to ATM VCIs below 32.
```

CAP_NET_BROADCAST

CODE **CAP_NET_BROADCAST**

```
CAP_NET_BROADCAST
```

```
    Allow broadcasting, listen to multicast.
```

CAP_NET_ADMIN

CODE **CAP_NET_ADMIN**

```
CAP_NET_ADMIN
```

```
    Allow interface configuration;
```

```
    Allow administration of IP firewall, masquerading and accounting;
```

```
    Allow setting debug option on sockets;
```

```
    Allow modification of routing tables;
```

```
    Allow setting arbitrary process / process group ownership on sockets;
```

```
    Allow binding to any address for transparent proxying;
```

```
    Allow setting TOS (type of service);
```

```
    Allow setting promiscuous mode;
```

```
    Allow clearing driver statistics;
```

```
    Allow multicasting;
```

```
    Allow read/write of device specific registers;
```

```
    Allow activation of ATM control sockets.
```

CAP_NET_RAW

CODE **CAP_NET_RAW**

```
CAP_NET_RAW
```

```
    Allow use of RAW sockets;
```

```
    Allow use of PACKET sockets.
```

CAP_IPC_LOCK

CODE **CAP_IPC_LOCK**

```
CAP_IPC_LOCK
```

```
    Allow locking of shared memory segments;
```

```
    Allow mlock and mlockall (which doesn't really have anything to do with IPC).
```

CAP_IPC_OWNER

CODE **CAP_IPC_OWNER**

```
CAP_IPC_OWNER
    Override IPC ownership checks.
```

CAP_SYS_MODULE

CODE CAP_SYS_MODULE

```
CAP_SYS_MODULE
    Insert and remove kernel modules    modify kernel without limit;
    Modify cap_bset.
```

CAP_SYS_RAWIO

CODE CAP_SYS_RAWIO

```
CAP_SYS_RAWIO
    Allow ioperm/iopl access;
    Allow sending USB messages to any device via /proc/bus/usb.
```

CAP_SYS_CHROOT

CODE CAP_SYS_CHROOT

```
CAP_SYS_CHROOT
    Allow use of chroot().
```

CAP_SYS_PTRACE

CODE CAP_SYS_PTRACE

```
CAP_SYS_PTRACE
    Allow ptrace() of any process.
```

CAP_SYS_PACCT

CODE CAP_SYS_PACCT

```
CAP_SYS_PACCT
    Allow configuration of process accounting.
```

CAP_SYS_ADMIN

CODE CAP_SYS_ADMIN

CAP_SYS_ADMIN

```
    Allow configuration of the secure attention key;
    Allow administration of the random device;
    Allow examination and configuration of disk quotas;
    Allow configuring the kernel's syslog (printk behaviour);
    Allow setting the domainname;
    Allow setting the hostname;
    Allow calling bdflush();
    Allow mount() and umount(), setting up new smb connection;
    Allow some autofs root ioctls;
    Allow nfsservctl; Allow VM86_REQUEST_IRQ;
    Allow to read/write pci config on alpha; Allow irix_prctl on mips (setstacksize);

    Allow flushing all cache on m68k (sys_cacheflush);
    Allow removing semaphores; Used instead of CAP_CHOWN to "chown" IPC message queues, semaphores and shared memory;
    Allow locking/unlocking of shared memory segment;
    Allow turning swap on/off;
    Allow forged pids on socket credentials passing;
    Allow setting readahead and flushing buffers on block devices;
    Allow setting geometry in floppy driver;
    Allow turning DMA on/off in xd driver;
    Allow administration of md devices (mostly the above, but some extra ioctls);
    Allow tuning the ide driver;
    Allow access to the nvram device;
    Allow administration of apm_bios, serial and bttv (TV) device;
    Allow manufacturer commands in isdn CAPI support driver;
    Allow reading nonstandardized portions of pci configuration space;
    Allow DDI debug ioctl on sbpcd driver;
    Allow setting up serial ports;
    Allow sending raw qic117 commands;
    Allow enabling/disabling tagged queuing on SCSI controllers and sending arbitrary SCSI commands;
    Allow setting encryption key on loopback filesystem.
```

CAP_SYS_BOOT

CODE CAP_SYS_BOOT

```
CAP_SYS_BOOT
    Allow use of reboot().
```

CAP_SYS_NICE

CODE CAP_SYS_NICE

```
CAP_SYS_NICE
    Allow raising priority and setting priority on other (different UID) processes;
CAP_SYS_RESOURCE
    Allow use of FIFO and roundrobin (realtime) scheduling on own processes and setting
    the scheduling algorithm used by another process.
```

CAP_SYS_RESOURCE

CODE CAP_SYS_RESOURCE

```
CAP_SYS_RESOURCE
    Override resource limits. Set resource limits;
    Override quota limits;
    Override reserved space on ext2 filesystem;
    Modify data journaling mode on ext3 filesystem
    (uses journaling resources); NOTE: ext2 honors fsuid when checking for
    resource overrides, so you can override using fsuid too;
    Override size restrictions on IPC message queues;
    Allow more than 64hz interrupts from the realtime clock;
    Override max number of consoles on console allocation;
    Override max number of keymaps.
```

CAP_SYS_TIME

CODE CAP_SYS_TIME

```
CAP_SYS_TIME
    Allow manipulation of system clock;
    Allow irix_stime on mips;
    Allow setting the realtime clock.
```

CAP_SYS_TTY_CONFIG

CODE CAP_SYS_TTY_CONFIG

```
CAP_SYS_TTY_CONFIG
    Allow configuration of tty devices; Allow vhangup() of tty.
```

CAP_MKNOD

CODE CAP_MKNOD

```
CAP_MKNOD
    Allow the privileged aspects of mknod().
```

CAP_LEASE

CODE CAP_LEASE

CAP_LEASE

Allow taking of leases on files.

This page is based on a document formerly found on our main website gentoo.org (<https://www.gentoo.org/>).

The following people contributed to the original document: **Ned Ludd, Adam Mondl**

They are listed here because wiki history does not allow for any external attribution. If you edit the wiki article, please do **not** add yourself here; your contributions are recorded on each article's associated history page.

Retrieved from "https://wiki.gentoo.org/index.php?title=Hardened/Overview_of_POSIX_capabilities&oldid=330422 (https://wiki.gentoo.org/index.php?title=Hardened/Overview_of_POSIX_capabilities&oldid=330422)"

- This page was last edited on 27 June 2015, at 17:55.
- [Privacy policy \(/wiki/Gentoo_Wiki:Privacy_policy\)](/wiki/Gentoo_Wiki:Privacy_policy)
- [About Gentoo Wiki \(/wiki/Gentoo_Wiki:About\)](/wiki/Gentoo_Wiki:About)
- [Disclaimers \(/wiki/Gentoo_Wiki:General_disclaimer\)](/wiki/Gentoo_Wiki:General_disclaimer)

© 2001–2022 Gentoo Foundation, Inc.

Gentoo is a trademark of the Gentoo Foundation, Inc. The contents of this document, unless otherwise expressly stated, are licensed under the CC-BY-SA-3.0 (<https://creativecommons.org/licenses/by-sa/3.0/>) license. The Gentoo Name and Logo Usage Guidelines (<https://www.gentoo.org/inside-gentoo/foundation/name-logo-guidelines.html>) apply.
